

OpenCV

Prerequisite: Before starting this exercise, you should make yourself familiar with Python and some necessary library, e.g., numpy, matplotlib, etc. One good tutorial can be found [here](#).

In this exercise you will:

- Learn about some basic image processing operations with OpenCV.
- Re-implement some basic image processing operations. This will help you to
- Have better understand about the image processing operations.
- Practice Python programming with Numpy library.

```
import cv2
import numpy as np
import sys
import matplotlib
from matplotlib import pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in
# the notebook
# rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of
plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

def rel_error(out, correct_out):
    return np.sum(abs(out.astype(np.float32) -
                        correct_out.astype(np.float32)) /
                    (abs(out.astype(np.float32)) +
                     abs(correct_out.astype(np.float32))))

# Checking OpenCV version
cv2.__version__

'4.5.5'
```

NOTICE:

In this lab exercise, we recommend to use OpenCV 3.x version, the documentations for OpenCV API can be found [here](#).

Load images

Use the function `cv2.imread()` to read an image. The image should be in the working directory or a full path of image should be given. The function will return a numpy matrix.

Second argument is a flag which specifies the way image should be read.

- `cv2.IMREAD_COLOR` - (1): Loads a color image. Any transparency (alpha channel) of image will be neglected. It is the **default flag**.
- `cv2.IMREAD_GRAYSCALE` - (0): Loads image in grayscale mode
- `cv2.IMREAD_UNCHANGED` - (-1): Loads image as such including alpha channel, if included.


NOTE: Color image loaded by OpenCV is in *Blue-Green-Red (BGR)* mode. But Matplotlib displays in *RGB* mode. So color images will not be displayed correctly in Matplotlib if image is read with OpenCV. We will discuss how to handle to display properly later.

```
img_gray = cv2.imread('imgs/opencv_logo.png', 0)

plt.figure(figsize=(20,10))
plt.subplot(131),
plt.imshow(img_gray, cmap='gray') # include cmap='gray' to display
gray image
plt.title('Gray'),plt.xticks([]), plt.yticks([])

img_color1= cv2.imread('imgs/opencv_logo.png', 1)
plt.subplot(132),plt.imshow(img_color1),
plt.title('Color'),plt.xticks([]), plt.yticks([])

img_color2= cv2.imread('imgs/opencv_logo.png',-1)
plt.subplot(133),plt.imshow(img_color2),
plt.title('Color'),plt.xticks([]), plt.yticks([])
plt.show()
```



Question: How many channels for each image: `img_gray`, `img_color1`, `img_color2`?

Your answer:

- `img_gray`:
- `img_color1`:
- `img_color2`:

Transformations

Scaling

Resize image using the function `cv2.resize`.

```
# Get list of available flags
flags = [i for i in dir(cv2) if i.startswith('INTER_')]
print (flags)

['INTER_AREA', 'INTER_BITS', 'INTER_BITS2', 'INTER_CUBIC',
 'INTER_LANCZOS4', 'INTER_LINEAR', 'INTER_LINEAR_EXACT', 'INTER_MAX',
 'INTER_NEAREST', 'INTER_NEAREST_EXACT', 'INTER_TAB_SIZE',
 'INTER_TAB_SIZE2']

img = cv2.imread('imgs/opencv_logo1.png', 1)
res = cv2.resize(img, None, fx=2.0, fy=2.0, interpolation =
cv2.INTER_CUBIC)
#OR
height, width = img.shape[:2]
res = cv2.resize(img, (2*width, 2*height), interpolation =
cv2.INTER_CUBIC)

#####
#####
# TO DO: Check the size of 'img' and 'res'?
#####
#####
print("img:", img.shape)
print("res:", res.shape)
#####
#####
#
                                END OF YOUR CODE
#
#####
#####

#####
#####
# TO DO: Resize 'img' so as to the smaller side is 500, while keeping
image
# ration unchanged.
#####
#####
shortest = min(height, width)
desired_shortest = 500
scale = desired_shortest / shortest
res = cv2.resize(img, None, fx=scale, fy=scale,
interpolation=cv2.INTER_CUBIC)
print("scaled img:", res.shape)
```

```
#####
#####
#
#
#
#####
#####

img: (378, 428, 3)
res: (756, 856, 3)
scaled img: (500, 566, 3)
```

Translation

Translation is the shifting of object's location. If you know the shift in (x, y) direction, let it be (t_x, t_y) , you can create the transformation matrix M as follows:

$$M = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix}$$

You can take make it into a Numpy array of type **np.float32** and pass it into `cv2.warpAffine()` function.

```
img = cv2.imread('imgs/opencv_logo1.png', 1)
rows,cols,_ = img.shape
M = np.float32([[1,0,100],[0,1,50]]) # Shift right by 100 and down by 50
dst = cv2.warpAffine(img,M,(cols,rows))

#####
#####
# TO DO: Observed that the bottom right of 'dst' image is lost.
Modifying the
# following codeline so as to the 'res' image is fully shown.
#####
#####
res = cv2.warpAffine(img,M,(cols+ 100,rows + 50))
#####
#####
#
#
#
#####
#####

plt.figure(figsize=(20,10))
plt.subplot(131),plt.imshow(img),
plt.title('Original'),plt.xticks([], plt.yticks([]))
plt.subplot(132),plt.imshow(dst),
plt.title('Shifted images'),plt.xticks([], plt.yticks([]))
plt.subplot(133),plt.imshow(res),
```

```
plt.title('Shifted images'),plt.xticks([]), plt.yticks([])
plt.show()
```



Rotation

Calculates an affine matrix of 2D rotation using `cv2.getRotationMatrix2D()`.

- 1st argument: center
- 2nd argument: angle (in degree)
- 3rd argument: scale

```
img = cv2.imread('imgs/opencv_logo1.png', 1)
H,W,_ = img.shape
#####
#####
# TO DO: Run the code to observe the output image.
# Modifying the code below so as to the 'dst' image has no black
padding.
#####
#####
M = cv2.getRotationMatrix2D((W/2,H/2),90,1)
dst = cv2.warpAffine(img,M,(W,H))

y, x = np.nonzero(dst)[:2]
dst = dst[np.min(y):np.max(y), np.min(x):np.max(x)]
#####
#####
#
#
#
#####
#####

plt.imshow(dst),
plt.title('Rotated images'),plt.xticks([]), plt.yticks([])
plt.show()
```



Changing color space - Grayscale

Grayscale values is converted from RGB values by a weighted sum of the R, G, and B components:

$$0.2989 \times R + 0.5870 \times G + 0.1140 \times B$$

```
# Split channels
```

```
img = cv2.imread('imgs/balls.jpg', 1)
```

```
plt.figure(figsize=(20,10))
```

```
plt.subplot(131),plt.imshow(img[:, :, 0], cmap='gray'),
```

```
plt.title('Blue channel'),plt.xticks([], plt.yticks([])
```


```
plt.subplot(132),plt.imshow(img[:, :, 1], cmap='gray'),
```

Run the following code section to compare your implementation of the `rgb2gray` function with OpenCV built-in function `cv2.cvtColor`.

```
plt.figure(figsize=(20,10))
```

```
plt.subplot(121),plt.imshow(img_gray1, cmap='gray'),
plt.title('My rgb2gray'),plt.xticks([]), plt.yticks([])
plt.subplot(122),plt.imshow(img_gray1 - img_gray2, cmap='gray'),
plt.title('Difference'),plt.xticks([]), plt.yticks([])
plt.show()

# Check your output: count
print('Testing rgb2gray')
print('Number of difference pixel is %d' % np.count_nonzero(img_gray1
- img_gray2))
```



Testing rgb2gray
Number of difference pixel is 81

Question: Does your implementation of rgb2gray function give the result that is exactly the same as OpenCV built-in function? Why?

Your answer: No. OpenCV uses a more accurate mathematical model as compared to the one listed above.

Changing color space - Detect object by color.

By converting BGR image to HSV, we can use this to extract a colored object. In HSV, it is more easier to represent a color than RGB color-space. In this exercise, we will try to extract blue, red, and yellow colored objects. So here is the method:

- Take each frame of the video
- Convert from BGR to HSV color-space
- We threshold the HSV image for a range of blue color
- Now extract the blue object alone, we can do whatever on that image we want.


```
# Get list of available flags
```

```
flags = [i for i in dir(cv2) if i.startswith('COLOR_')]  
print(flags)
```

```
['COLOR_BAYER_BG2BGR', 'COLOR_BAYER_BG2BGRA', 'COLOR_BAYER_BG2BGR_EA',  
'COLOR_BAYER_BG2BGR_VNG', 'COLOR_BAYER_BG2GRAY', 'COLOR_BAYER_BG2RGB',  
'COLOR_BAYER_BG2RGBA', 'COLOR_BAYER_BG2RGB_EA',  
'COLOR_BAYER_BG2RGB_VNG', 'COLOR_BAYER_BGGR2BGR',  
'COLOR_BAYER_BGGR2BGRA', 'COLOR_BAYER_BGGR2BGR_EA',  
'COLOR_BAYER_BGGR2BGR_VNG', 'COLOR_BAYER_BGGR2GRAY',  
'COLOR_BAYER_BGGR2RGB', 'COLOR_BAYER_BGGR2RGBA',  
'COLOR_BAYER_BGGR2RGB_EA', 'COLOR_BAYER_BGGR2RGB_VNG',  
'COLOR_BAYER_GB2BGR', 'COLOR_BAYER_GB2BGRA', 'COLOR_BAYER_GB2BGR_EA',  
'COLOR_BAYER_GB2BGR_VNG', 'COLOR_BAYER_GB2GRAY', 'COLOR_BAYER_GB2RGB',  
'COLOR_BAYER_GB2RGBA', 'COLOR_BAYER_GB2RGB_EA',  
'COLOR_BAYER_GB2RGB_VNG', 'COLOR_BAYER_GBRG2BGR',  
'COLOR_BAYER_GBRG2BGRA', 'COLOR_BAYER_GBRG2BGR_EA',  
'COLOR_BAYER_GBRG2BGR_VNG', 'COLOR_BAYER_GBRG2GRAY',  
'COLOR_BAYER_GBRG2RGB', 'COLOR_BAYER_GBRG2RGBA',  
'COLOR_BAYER_GBRG2RGB_EA', 'COLOR_BAYER_GBRG2RGB_VNG',  
'COLOR_BAYER_GR2BGR', 'COLOR_BAYER_GR2BGRA', 'COLOR_BAYER_GR2BGR_EA',  
'COLOR_BAYER_GR2BGR_VNG', 'COLOR_BAYER_GR2GRAY', 'COLOR_BAYER_GR2RGB',  
'COLOR_BAYER_GR2RGBA', 'COLOR_BAYER_GR2RGB_EA',  
'COLOR_BAYER_GR2RGB_VNG', 'COLOR_BAYER_GRBG2BGR',  
'COLOR_BAYER_GRBG2BGRA', 'COLOR_BAYER_GRBG2BGR_EA',  
'COLOR_BAYER_GRBG2BGR_VNG', 'COLOR_BAYER_GRBG2GRAY',  
'COLOR_BAYER_GRBG2RGB', 'COLOR_BAYER_GRBG2RGBA',  
'COLOR_BAYER_GRBG2RGB_EA', 'COLOR_BAYER_GRBG2RGB_VNG',  
'COLOR_BAYER_RG2BGR', 'COLOR_BAYER_RG2BGRA', 'COLOR_BAYER_RG2BGR_EA',  
'COLOR_BAYER_RG2BGR_VNG', 'COLOR_BAYER_RG2GRAY', 'COLOR_BAYER_RG2RGB',  
'COLOR_BAYER_RG2RGBA', 'COLOR_BAYER_RG2RGB_EA',  
'COLOR_BAYER_RG2RGB_VNG', 'COLOR_BAYER_RGGB2BGR',  
'COLOR_BAYER_RGGB2BGRA', 'COLOR_BAYER_RGGB2BGR_EA',  
'COLOR_BAYER_RGGB2BGR_VNG', 'COLOR_BAYER_RGGB2GRAY',  
'COLOR_BAYER_RGGB2RGB', 'COLOR_BAYER_RGGB2RGBA',  
'COLOR_BAYER_RGGB2RGB_EA', 'COLOR_BAYER_RGGB2RGB_VNG',  
'COLOR_BGR2BGR555', 'COLOR_BGR2BGR565', 'COLOR_BGR2BGRA',  
'COLOR_BGR2GRAY', 'COLOR_BGR2HLS', 'COLOR_BGR2HLS_FULL',  
'COLOR_BGR2HSV', 'COLOR_BGR2HSV_FULL', 'COLOR_BGR2LAB',  
'COLOR_BGR2LUV', 'COLOR_BGR2Lab', 'COLOR_BGR2Luv', 'COLOR_BGR2RGB',  
'COLOR_BGR2RGBA', 'COLOR_BGR2XYZ', 'COLOR_BGR2YCrCb',  
'COLOR_BGR2YCrCb', 'COLOR_BGR2YUV', 'COLOR_BGR2YUV_I420',  
'COLOR_BGR2YUV_IYUV', 'COLOR_BGR2YUV_YV12', 'COLOR_BGR5552BGR',  
'COLOR_BGR5552BGRA', 'COLOR_BGR5552GRAY', 'COLOR_BGR5552RGB',  
'COLOR_BGR5552RGBA', 'COLOR_BGR5652BGR', 'COLOR_BGR5652BGRA',  
'COLOR_BGR5652GRAY', 'COLOR_BGR5652RGB', 'COLOR_BGR5652RGBA',  
'COLOR_BGRA2BGR', 'COLOR_BGRA2BGR555', 'COLOR_BGRA2BGR565',  
'COLOR_BGRA2GRAY', 'COLOR_BGRA2RGB', 'COLOR_BGRA2RGBA',  
'COLOR_BGRA2YUV_I420', 'COLOR_BGRA2YUV_IYUV', 'COLOR_BGRA2YUV_YV12',  
'COLOR_BayerBG2BGR', 'COLOR_BayerBG2BGRA', 'COLOR_BayerBG2BGR_EA',
```

'COLOR_BayerBG2BGR_VNG', 'COLOR_BayerBG2GRAY', 'COLOR_BayerBG2RGB',
'COLOR_BayerBG2RGBA', 'COLOR_BayerBG2RGB_EA', 'COLOR_BayerBG2RGB_VNG',
'COLOR_BayerBGGR2BGR', 'COLOR_BayerBGGR2BGRA',
'COLOR_BayerBGGR2BGR_EA', 'COLOR_BayerBGGR2BGR_VNG',
'COLOR_BayerBGGR2GRAY', 'COLOR_BayerBGGR2RGB', 'COLOR_BayerBGGR2RGBA',
'COLOR_BayerBGGR2RGB_EA', 'COLOR_BayerBGGR2RGB_VNG',
'COLOR_BayerGB2BGR', 'COLOR_BayerGB2BGRA', 'COLOR_BayerGB2BGR_EA',
'COLOR_BayerGB2BGR_VNG', 'COLOR_BayerGB2GRAY', 'COLOR_BayerGB2RGB',
'COLOR_BayerGB2RGBA', 'COLOR_BayerGB2RGB_EA', 'COLOR_BayerGB2RGB_VNG',
'COLOR_BayerGBRG2BGR', 'COLOR_BayerGBRG2BGRA',
'COLOR_BayerGBRG2BGR_EA', 'COLOR_BayerGBRG2BGR_VNG',
'COLOR_BayerGBRG2GRAY', 'COLOR_BayerGBRG2RGB', 'COLOR_BayerGBRG2RGBA',
'COLOR_BayerGBRG2RGB_EA', 'COLOR_BayerGBRG2RGB_VNG',
'COLOR_BayerGR2BGR', 'COLOR_BayerGR2BGRA', 'COLOR_BayerGR2BGR_EA',
'COLOR_BayerGR2BGR_VNG', 'COLOR_BayerGR2GRAY', 'COLOR_BayerGR2RGB',
'COLOR_BayerGR2RGBA', 'COLOR_BayerGR2RGB_EA', 'COLOR_BayerGR2RGB_VNG',
'COLOR_BayerGRBG2BGR', 'COLOR_BayerGRBG2BGRA',
'COLOR_BayerGRBG2BGR_EA', 'COLOR_BayerGRBG2BGR_VNG',
'COLOR_BayerGRBG2GRAY', 'COLOR_BayerGRBG2RGB', 'COLOR_BayerGRBG2RGBA',
'COLOR_BayerGRBG2RGB_EA', 'COLOR_BayerGRBG2RGB_VNG',
'COLOR_BayerRG2BGR', 'COLOR_BayerRG2BGRA', 'COLOR_BayerRG2BGR_EA',
'COLOR_BayerRG2BGR_VNG', 'COLOR_BayerRG2GRAY', 'COLOR_BayerRG2RGB',
'COLOR_BayerRG2RGBA', 'COLOR_BayerRG2RGB_EA', 'COLOR_BayerRG2RGB_VNG',
'COLOR_BayerRGG2BGR', 'COLOR_BayerRGG2BGRA',
'COLOR_BayerRGG2BGR_EA', 'COLOR_BayerRGG2BGR_VNG',
'COLOR_BayerRGG2GRAY', 'COLOR_BayerRGG2RGB', 'COLOR_BayerRGG2RGBA',
'COLOR_BayerRGG2RGB_EA', 'COLOR_BayerRGG2RGB_VNG',
'COLOR_COLORCVT_MAX', 'COLOR_GRAY2BGR', 'COLOR_GRAY2BGR555',
'COLOR_GRAY2BGR565', 'COLOR_GRAY2BGRA', 'COLOR_GRAY2RGB',
'COLOR_GRAY2RGBA', 'COLOR_HLS2BGR', 'COLOR_HLS2BGR_FULL',
'COLOR_HLS2RGB', 'COLOR_HLS2RGB_FULL', 'COLOR_HSV2BGR',
'COLOR_HSV2BGR_FULL', 'COLOR_HSV2RGB', 'COLOR_HSV2RGB_FULL',
'COLOR_LAB2BGR', 'COLOR_LAB2LBGR', 'COLOR_LAB2LRGB', 'COLOR_LAB2RGB',
'COLOR_LBGR2LAB', 'COLOR_LBGR2LUV', 'COLOR_LBGR2Lab',
'COLOR_LBGR2Luv', 'COLOR_LRGB2LAB', 'COLOR_LRGB2LUV',
'COLOR_LRGB2Lab', 'COLOR_LRGB2Luv', 'COLOR_LUV2BGR', 'COLOR_LUV2LBGR',
'COLOR_LUV2LRGB', 'COLOR_LUV2RGB', 'COLOR_Lab2BGR', 'COLOR_Lab2LBGR',
'COLOR_Lab2LRGB', 'COLOR_Lab2RGB', 'COLOR_Luv2BGR', 'COLOR_Luv2LBGR',
'COLOR_Luv2LRGB', 'COLOR_Luv2RGB', 'COLOR_M_RGBA2RGBA',
'COLOR_RGB2BGR', 'COLOR_RGB2BGR555', 'COLOR_RGB2BGR565',
'COLOR_RGB2BGRA', 'COLOR_RGB2GRAY', 'COLOR_RGB2HLS',
'COLOR_RGB2HLS_FULL', 'COLOR_RGB2HSV', 'COLOR_RGB2HSV_FULL',
'COLOR_RGB2LAB', 'COLOR_RGB2LUV', 'COLOR_RGB2Lab', 'COLOR_RGB2Luv',
'COLOR_RGB2RGBA', 'COLOR_RGB2XYZ', 'COLOR_RGB2YCR_CB',
'COLOR_RGB2YCrCb', 'COLOR_RGB2YUV', 'COLOR_RGB2YUV_I420',
'COLOR_RGB2YUV_IYUV', 'COLOR_RGB2YUV_YV12', 'COLOR_RGBA2BGR',
'COLOR_RGBA2BGR555', 'COLOR_RGBA2BGR565', 'COLOR_RGBA2BGRA',
'COLOR_RGBA2GRAY', 'COLOR_RGBA2M_RGBA', 'COLOR_RGBA2RGB',
'COLOR_RGBA2YUV_I420', 'COLOR_RGBA2YUV_IYUV', 'COLOR_RGBA2YUV_YV12',
'COLOR_RGBA2mRGBA', 'COLOR_XYZ2BGR', 'COLOR_XYZ2RGB',

```

'COLOR_YCR_CB2BGR', 'COLOR_YCR_CB2RGB', 'COLOR_YCrCb2BGR',
'COLOR_YCrCb2RGB', 'COLOR_YUV2BGR', 'COLOR_YUV2BGRA_I420',
'COLOR_YUV2BGRA_IYUV', 'COLOR_YUV2BGRA_NV12', 'COLOR_YUV2BGRA_NV21',
'COLOR_YUV2BGRA_UYNV', 'COLOR_YUV2BGRA_UYVY', 'COLOR_YUV2BGRA_Y422',
'COLOR_YUV2BGRA_YUNV', 'COLOR_YUV2BGRA_YUY2', 'COLOR_YUV2BGRA_YUYV',
'COLOR_YUV2BGRA_YV12', 'COLOR_YUV2BGRA_YVYU', 'COLOR_YUV2BGR_I420',
'COLOR_YUV2BGR_IYUV', 'COLOR_YUV2BGR_NV12', 'COLOR_YUV2BGR_NV21',
'COLOR_YUV2BGR_UYNV', 'COLOR_YUV2BGR_UYVY', 'COLOR_YUV2BGR_Y422',
'COLOR_YUV2BGR_YUNV', 'COLOR_YUV2BGR_YUY2', 'COLOR_YUV2BGR_YUYV',
'COLOR_YUV2BGR_YV12', 'COLOR_YUV2BGR_YVYU', 'COLOR_YUV2GRAY_420',
'COLOR_YUV2GRAY_I420', 'COLOR_YUV2GRAY_IYUV', 'COLOR_YUV2GRAY_NV12',
'COLOR_YUV2GRAY_NV21', 'COLOR_YUV2GRAY_UYNV', 'COLOR_YUV2GRAY_UYVY',
'COLOR_YUV2GRAY_Y422', 'COLOR_YUV2GRAY_YUNV', 'COLOR_YUV2GRAY_YUY2',
'COLOR_YUV2GRAY_YUYV', 'COLOR_YUV2GRAY_YV12', 'COLOR_YUV2GRAY_YVYU',
'COLOR_YUV2RGB', 'COLOR_YUV2RGBA_I420', 'COLOR_YUV2RGBA_IYUV',
'COLOR_YUV2RGBA_NV12', 'COLOR_YUV2RGBA_NV21', 'COLOR_YUV2RGBA_UYNV',
'COLOR_YUV2RGBA_UYVY', 'COLOR_YUV2RGBA_Y422', 'COLOR_YUV2RGBA_YUNV',
'COLOR_YUV2RGBA_YUY2', 'COLOR_YUV2RGBA_YUYV', 'COLOR_YUV2RGBA_YV12',
'COLOR_YUV2RGBA_YVYU', 'COLOR_YUV2RGB_I420', 'COLOR_YUV2RGB_IYUV',
'COLOR_YUV2RGB_NV12', 'COLOR_YUV2RGB_NV21', 'COLOR_YUV2RGB_UYNV',
'COLOR_YUV2RGB_UYVY', 'COLOR_YUV2RGB_Y422', 'COLOR_YUV2RGB_YUNV',
'COLOR_YUV2RGB_YUY2', 'COLOR_YUV2RGB_YUYV', 'COLOR_YUV2RGB_YV12',
'COLOR_YUV2RGB_YVYU', 'COLOR_YUV420P2BGR', 'COLOR_YUV420P2BGRA',
'COLOR_YUV420P2GRAY', 'COLOR_YUV420P2RGB', 'COLOR_YUV420P2RGBA',
'COLOR_YUV420SP2BGR', 'COLOR_YUV420SP2BGRA', 'COLOR_YUV420SP2GRAY',
'COLOR_YUV420SP2RGB', 'COLOR_YUV420SP2RGBA', 'COLOR_YUV420p2BGR',
'COLOR_YUV420p2BGRA', 'COLOR_YUV420p2GRAY', 'COLOR_YUV420p2RGB',
'COLOR_YUV420p2RGBA', 'COLOR_YUV420sp2BGR', 'COLOR_YUV420sp2BGRA',
'COLOR_YUV420sp2GRAY', 'COLOR_YUV420sp2RGB', 'COLOR_YUV420sp2RGBA',
'COLOR_mRGBA2RGBA']

```

```
frame = cv2.imread('imgs/balls.jpg', 1)
```

```
# Convert BGR to RGB, now you will see the color of 'frame' image
# is displayed properly.
```

```
frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
```

```
# Convert BGR to HSV
```

```
hsv = cv2.cvtColor(frame, cv2.COLOR_RGB2HSV)
```

```
# define range of blue color in HSV
```

```
lower_blue = np.array([110,50,50])
```

```
upper_blue = np.array([130,255,255])
```

```
# Threshold the HSV image to get only blue colors
```

```
mask = cv2.inRange(hsv, lower_blue, upper_blue)
```


```
# Bitwise-AND mask and original image
```

```
res = cv2.bitwise_and(frame,frame, mask= mask)
```

```
#####
#####
# TO DO: Implement masks for red and yellow balls.
#####
#####
# Lower red hue
lower_red1 = np.array([0, 50, 50])
upper_red1 = np.array([20, 255, 255])
red_lower_mask = cv2.inRange(hsv, lower_red1, upper_red1)
# Upper red hue
lower_red2 = np.array([170, 50, 50])
upper_red2 = np.array([180, 255, 255])
red_upper_mask= cv2.inRange(hsv, lower_red2, upper_red2)
#total red
red_mask = red_lower_mask + red_upper_mask
red_res = cv2.bitwise_and(frame,frame, mask= red_mask)

# Yellow hue
lower_yellow = np.array([20, 50, 50])
upper_yellow = np.array([40, 255, 255])
yellow_mask= cv2.inRange(hsv, lower_yellow, upper_yellow)
yellow_res = cv2.bitwise_and(frame, frame, mask=yellow_mask)
#####
#####
#
#
#
#####
#####

plt.figure(figsize=(20,10))
plt.subplot(131),plt.imshow(frame),
plt.title('Original'),plt.xticks([]), plt.yticks([])
plt.subplot(132),plt.imshow(mask, cmap='gray'),
plt.title('Mask'),plt.xticks([]), plt.yticks([])
plt.subplot(133),plt.imshow(res),
#plt.subplot(133),plt.imshow(red_res)
#plt.subplot(133),plt.imshow(yellow_res)
plt.title('Output'),plt.xticks([]), plt.yticks([])
plt.show()
```



2D Convolution (Image Filtering)

OpenCV provides a function, `cv2.filter2D`, to convolve a kernel with an image.

```
def convolution_naive(x, F, conv_param):
    """
    A naive implementation of a convolutional filter.

    The input consists of a gray scale image x (1 channel) with height
    H and width
    W. We convolve each input with filter F, which has height HH and
    width HH.

    Input:
    - x: Input data of shape (H, W)
    - F: Filter weights of shape (HH, WW)
    - conv_param: A dictionary with the following keys:
        - 'stride': The number of pixels between adjacent receptive
        fields in the
        horizontal and vertical directions.
        - 'pad': The number of pixels that will be used to zero-pad the
        input.

    Return:
    - out: Output data, of shape (H', W') where H' and W' are given by
         $H' = 1 + (H + 2 * pad - HH) / stride$ 
         $W' = 1 + (W + 2 * pad - WW) / stride$ 
    """

    stride = conv_param['stride']
    pad = conv_param['pad']
    H, W = x.shape
    HH, WW = F.shape
    H_prime = int(1 + (H + 2 * pad - HH) / stride)
    W_prime = int(1 + (W + 2 * pad - WW) / stride)
    x_pad = np.lib.pad(x, ((pad, pad), (pad, pad)),\
        'constant', constant_values=(0))
    out = np.zeros((H_prime, W_prime), dtype=x.dtype)
    print(x_pad.shape)

#####
#####
    # TODO: Implement the convolutional forward pass.
    #
    # Hint: Using 2 nested for-loop to calculate each pixel of the
    output image.#

#####
#####
    for i in range(0, W_prime, stride):
```

```

        for j in range(0, H_prime, stride):
            subset = x_pad[i: i + WW, j: j + HH]
            out[i, j] = (subset * F).sum()

#####
#####
#                                     END OF YOUR CODE
#

#####
#####
return out

```

Run the following code section to test your implementation of the convolution_naive function

```

x_shape = (5, 5)
F_shape = (3, 3)
x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
F = np.linspace(-0.2, 0.3, num=np.prod(F_shape)).reshape(F_shape)
conv_param = {'stride': 1, 'pad': 1}

out = convolution_naive(x, F, conv_param)
correct_out = np.array( [[ 0.0075,      0.030625,   0.0521875,
0.07375,      0.0475   ],
                        [ 0.114375,   0.1725,     0.18375,   0.195,
0.10875   ],
                        [ 0.1753125,  0.22875,     0.24,
0.25125,     0.1228125],
                        [ 0.23625,     0.285,       0.29625,   0.3075,
0.136875 ],
                        [ 0.0075,     -0.05375,   -0.0603125, -
0.066875,   -0.1025   ]])
# print(correct_out.shape)
# print(out)

# Compare your output to ours; difference should be very small
print('Testing convolution_naive')
print('difference: ', rel_error(out, correct_out))

(7, 7)
Testing convolution_naive
difference:  0.0

# List of available BORDER effect
flags = [i for i in dir(cv2) if i.startswith('BORDER_')]
print(flags)

['BORDER_CONSTANT', 'BORDER_DEFAULT', 'BORDER_ISOLATED',
'BORDER_REFLECT', 'BORDER_REFLECT101', 'BORDER_REFLECT_101',
'BORDER_REPLICATE', 'BORDER_TRANSPARENT', 'BORDER_WRAP']

```

Averaging filter

This is done by convolving image with a normalized box filter. A 5×5 normalized box filter would look like below:

$$K = \frac{1}{25} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

```
# Convert image data type from uint8 to float32.
img = cv2.imread('imgs/text.png', 1).astype(np.float32)
img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

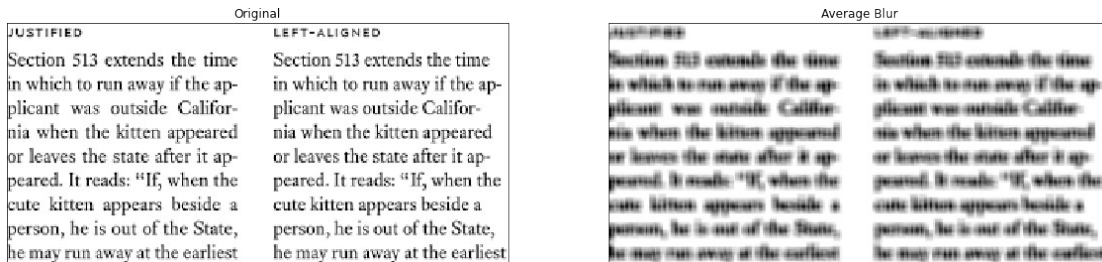
kernel = np.zeros((5,5), np.float32)
#####
#####
# TODO: Create a 5x5 kernel as K shown above.
#
#####
#####
kernel = 1/25 * np.ones((5,5), np.float32)
#####
#####
#
#                               END OF YOUR CODE
#
#####
#####
blur_2dfilter = cv2.filter2D(img,-1,kernel)

# The above codes can be replaced by the following code line.
blur = cv2.blur(img,(5,5))

# Check your output; difference should be around 4e-3
print('Testing convolution_naive')
print('difference: ', rel_error(blur_2dfilter, blur))

# Visualize the output image
plt.figure(figsize=(20,10))
plt.subplot(121),plt.imshow(img, cmap='gray'),
plt.title('Original'),plt.xticks([], plt.yticks([]))
plt.subplot(122),plt.imshow(blur, cmap='gray'),
plt.title('Average Blur'),plt.xticks([], plt.yticks([]))
plt.show()

Testing convolution_naive
difference:  0.0035056125
```



Gaussian Blurring

Here is the 1D Gaussian distribution:

$$G(x) = \frac{1}{\sigma \sqrt{2\pi}} \exp\left(-\frac{x^2}{\sigma^2}\right)$$

1D Gaussian

Similarly, we have 2D Gaussian distribution.

$$G(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{\sigma^2}\right)$$

The nearest neighboring pixels have the most influence. 2D Gaussian

```
img = cv2.imread('imgs/text.png', 1).astype(np.float32)/255.0
img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
```

```
gaussian_kernel_XY = np.zeros((5,5), np.float32)
#####
#####
# TODO: Create a 5x5 kernel, 'gaussian_kernel_XY', which approximates
# the
# Gaussian function with sigma=1.
# Hint: + You should NOT manually create the kernel.
#       + Use the 'cv2.getGaussianKernel' function to create 1D
#       + Use the associative property of convolution to create 2D
#       Gaussian. A
# useful reference:
# https://blogs.mathworks.com/steve/2006/10/04/separable-convolution/
#####
#####
gaussian = cv2.getGaussianKernel(5, 1)
gaussian_kernel_XY = gaussian.dot(gaussian.T)
#####
#####
#
#
#####
#
#
#####
#
#
#####
```



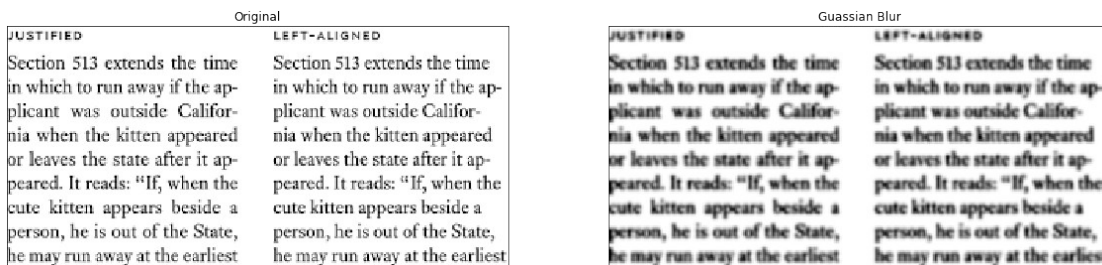
```
#####
blur_2dfilter = cv2.filter2D(img,-1,gaussian_kernel_XY)

# The above codes can be replaced by the following code line.
blur = cv2.GaussianBlur(img,(5,5),1)

# Check your output; difference should be around 4e-3
print('Testing convolution_naive')
print('difference: ', rel_error(blur_2dfilter, blur))

# Visualize the output image
plt.figure(figsize=(20,10))
plt.subplot(121),plt.imshow(img, cmap='gray'),
plt.title('Original'),plt.xticks([]), plt.yticks([])
plt.subplot(122),plt.imshow(blur, cmap='gray'),
plt.title('Guassain Blur'),plt.xticks([]), plt.yticks([])
plt.show()
```

Testing convolution_naive
difference: 0.0042602094



QUESTION: Provide your comments on the outputs of a *average filter* and a *Gaussian filter*? Which one is more preferable?

Your answer: Gaussian filter has a more smoother and natural blur as compared to the sharper edges from average filter

Median Filter

Example:

- **Odd** number of elements: $X = [2, 5, 1, 0, 9] \rightarrow X_{sorted} = [0, 1, 2, 5, 9] \Rightarrow \text{median} = 2$
- **Even** number of elements:
 - Option 1: $X = [5, 1, 0, 9] \rightarrow X_{sorted} = [0, 1, 5, 9] \Rightarrow \text{median} = 1$
 - Option 2: $X = [5, 1, 0, 9] \rightarrow X_{sorted} = [0, 1, 5, 9] \Rightarrow \text{median} = (1+5)/2 = 3$

Implement a function to find median value with 'option 1'.

```
def findMedian(x):
    out = 0
```

#####

```
#####  
# TODO: Implement the function to find median value of array x.  
#  
# NOTE: You should see that the `median` numpy built-in function  
is based #  
# on option 2.  
  
#####  
#####  
x = np.sort(np.array(x).flatten())  
out = x[len(x) // 2] if len(x) % 2 else x[(len(x) // 2) - 1]  
  
#####  
#####  
#  
#  
# END OF YOUR CODE  
  
#####  
#####  
return out  
  
print ('Numpy median: ', np.median([[5,1],[0,9]]))  
print ('Numpy median: ', np.median([2,5,1,0,9]))  
print ('findMedian: ', findMedian([[5,1],[0,9]]))  
print ('findMedian: ', findMedian([2,5,1,0,9]))  
  
Numpy median: 3.0  
Numpy median: 2.0  
findMedian: 1  
findMedian: 2  
  
img = cv2.imread('imgs/SaltAndPepperNoise.jpg', 0)  
median = cv2.medianBlur(img,5)  
gau_blur = cv2.GaussianBlur(img,(5,5),1)  
  
plt.figure(figsize=(20,10))  
plt.subplot(131),plt.imshow(img, 'gray')  
plt.title('Original'),plt.xticks([],plt.yticks([]))  
plt.subplot(132),plt.imshow(median, 'gray')  
plt.title('Median Blur'),plt.xticks([],plt.yticks([]))  
plt.subplot(133),plt.imshow(gau_blur, 'gray')  
plt.title('Gaussian Blur'),plt.xticks([],plt.yticks([]))  
plt.show()
```



```
#####
#####
    return out

img = cv2.imread('imgs/SaltAndPepperNoise.jpg', 0)
mymedian = myMedianBlur(img,5)
median = cv2.medianBlur(img,5)

# Note that your implementation is NOT necessary to provide
# the identical output as OpenCV built-in function. However,
# it should visually very similar.
plt.figure(figsize=(16,8))
plt.subplot(121),plt.imshow(median, 'gray')
plt.title('Opencv Median Blur'),plt.xticks([],plt.yticks([]))
plt.subplot(122),plt.imshow(median, 'gray')
plt.title('My Median Blur'),plt.xticks([],plt.yticks([]))
plt.show()
```



Image gradient

For 1-D continuous function $f(x)$, the gradient is given as:

$$D_x[f(x)] = \frac{d}{dx} f(x) = \lim_{\Delta x \rightarrow 0} \frac{f(x+\Delta x) - f(x)}{\Delta x}, \text{ or } \lim_{\Delta x \rightarrow 0} \frac{f(x+\Delta x) - f(x-\Delta x)}{2\Delta x}$$

For 1-D discrete function $f[n]$, the gradient becomes difference.

$$D_n[f[n]] = f[n+1] - f[n], \text{ or } \frac{f[n+1] - f[n-1]}{2}$$

The kernel to find gradient in 1-D discrete function is $[1, 0, -1]$.

```


img = cv2.imread('imgs/banded_vertical.jpg', 0).astype(np.float32)

#####
#####
# TODO: Create a 3x3 kernel, Kx, to find the gradient in x-axis of an
image.#
#####
#####
Kx = np.zeros((3, 3))
Kx[:, 0] = 1
Kx[:, 2] = -1
print(Kx)
#####
#####
#                                     END OF YOUR CODE
#
#####
#####
dstx = cv2.filter2D(img,-1, Kx)

plt.figure(figsize=(10,5))
plt.subplot(121),plt.imshow(img, cmap='gray')
plt.title('Original'),plt.xticks([]),plt.yticks([])
plt.subplot(122),plt.imshow(np.abs(dstx), cmap='gray')
plt.title('Output 1'),plt.xticks([]),plt.yticks([])
plt.show()

[[ 1.  0. -1.]
 [ 1.  0. -1.]
 [ 1.  0. -1.]]

```



```


img = cv2.imread('imgs/banded_horizontal.jpg', 0).astype(np.float32)

```

```
#####
#####
# TODO: Create a 3x3 kernel, Ky, to find the gradient in y-axis of an
image.#
#####
#####
Ky = np.zeros((3, 3))
Ky[0] = 1
Ky[2] = -1
print(Ky)
#####
#####
#
#
#
#####
#####
dsty = cv2.filter2D(img, -1, Ky)

plt.figure(figsize=(10,5))
plt.subplot(121),plt.imshow(img, 'gray')
plt.title('Original'),plt.xticks([]),plt.yticks([])
plt.subplot(122),plt.imshow(np.abs(dsty), 'gray')
plt.title('Output'),plt.xticks([]),plt.yticks([])
plt.show()

[[ 1.  1.  1.]
 [ 0.  0.  0.]
 [-1. -1. -1.]]
```



Question: What do the kernel Kx and Ky do in *image processing*?

Answer: Kx and Ky are convolution kernels that provides the horizontal and vertical image gradients respectively

Two directions:

- Find the difference: in the two directions:

$$g_x[m,n] = f[m+1,n] - f[m-1,n]$$

$$g_y[m,n] = f[m,n+1] - f[m,n-1]$$

- Find the magnitude and direction of the gradient vector:

$$\|g[m,n]\| = \sqrt{g_x^2[m,n] + g_y^2[m,n]}$$

$$\angle g[m,n] = \tan^{-1} \left(\frac{g_y[m,n]}{g_x[m,n]} \right)$$

```
img = cv2.imread('imgs/chequered.jpg', 0).astype(np.float32)
```


```
#####
#####
# TODO: Using the theory provided above, compute the magnitude of 2
#
# direction image gradient.
#
#####
#####
gx, gy = img.copy(), img.copy()
gx[:-2] -= img[2:]
gy[:, :-2] -= img[:, 2:]
dst1 = np.sqrt(gx ** 2 + gy ** 2)
#####
#####
#
#                               END OF YOUR CODE
#
#####
#####
```

You can achieve a similar (NOT identical) output with the following code line.

```
K = np.array([[0, 1, 0],
              [1, -4, 1],
              [0, 1, 0]], dtype=np.float32)
dst2 = cv2.filter2D(img, -1, K)
```

```
plt.figure(figsize=(20,10))
plt.subplot(131),plt.imshow(img, 'gray')
plt.title('Original'),plt.xticks([],plt.yticks([]))
plt.subplot(132),plt.imshow(np.abs(dst1), 'gray')
```

```
plt.title('Output 1'),plt.xticks([]),plt.yticks([])
plt.subplot(133),plt.imshow(np.abs(dst2), 'gray')
plt.title('Output 2'),plt.xticks([]),plt.yticks([])
plt.show()
```



Histogram

- It is a graphical representation of the intensity distribution of an image.
- It quantifies the number of pixels for each intensity value considered.

Histogram equalization

- Equalization implies mapping one distribution (the given histogram) to another distribution (a wider and more uniform distribution of intensity values) so the intensity values are spreaded over the whole range.
- To accomplish the equalization effect, the remapping should be the cumulative distribution function (cdf) (more details, refer to Learning OpenCV). For the histogram $H(i)$, its cumulative distribution $H'(i)$ is:

$$H'(i) = \sum_{0 \leq j < i} H(j)$$

- To use this as a remapping function, we have to normalize $H'(i)$ such that the maximum value is 255 (or the maximum value for the intensity of the image). From the example above, the cumulative function is:

cumulative distribution function

- Finally, we use a simple remapping procedure to obtain the intensity values of the equalized image:

$$equalized(x,y) = H'(src(x,y))$$

Histogram Equalization

```
img = cv2.imread('imgs/sudoku-original.jpg',0)
W,H = img.shape
img_eq = cv2.equalizeHist(img)
```

```
hist = np.histogram(img, bins=256, range=(0.0, 255.0))
hist_eq = np.histogram(img_eq, bins=256, range=(0.0, 255.0))
```



```

plt.figure(figsize=(10,15))
plt.subplot(321),plt.imshow(img, cmap='gray'),plt.title('Original
Image'),plt.xticks([]),plt.yticks([])
plt.subplot(322),plt.imshow(img_eq, cmap='gray'),plt.title('Equalized
Image'),plt.xticks([]),plt.yticks([])
plt.subplot(323),plt.hist(img.ravel(), bins=256, range=(0.0,
255.0)),plt.title('Original Image: Histogram')
plt.subplot(324),plt.hist(img_eq.ravel(), bins=256, range=(0.0,
255.0)),plt.title('Equalized Image: Histogram')
plt.subplot(325),plt.plot(range(0,256),np.cumsum(hist[0])*255/(W*H)),p
lt.title('Original Image: normalized CDF')
plt.subplot(326),plt.plot(range(0,256),np.cumsum(hist_eq[0])*255/(W*H)
),plt.title('Equalized Image: normalized CDF')
plt.show()

```


Original Image




Equalized Image




Original Image: Histogram




Equalized Image: Histogram



Original Image: normalized CDF



Equalized Image: normalized CDF



QUIZ: Is histogram equalization reversible?

Your answer: Histogram equalization is a many-to-one mapping, and therefore not reversible

```
def myEqualizeHist(img):
    """
    A implementation of a histogram equalization for image of `uint8`
    data type.
    """
    out = img

#####
#####
    # TODO: Implement the histogram equalization function.
    #

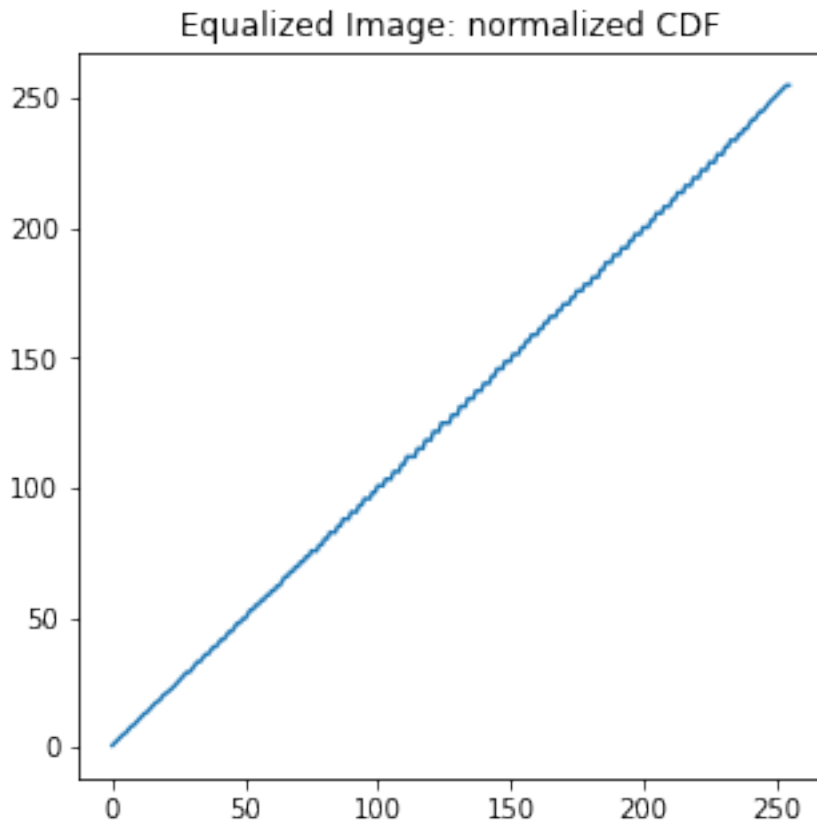
#####
#####
    freq, bins = np.histogram(img.flatten(), 256, [0, 256])
    cdf = np.cumsum(freq).astype(float)
    # mask all pixels with value 0 and replace them with the mean of
    the pixel values
    cdf_m = np.ma.masked_equal(cdf,0)
    cdf_m = (cdf_m - cdf_m.min())*255/(cdf_m.max()-cdf_m.min())
    cdf_final = np.ma.filled(cdf_m,0).astype('uint8')
    out = cdf_final[img]

#####
#####
    #                                     END OF YOUR CODE
    #

#####
#####
    return out

# Verify the correctness of your implementation by plotting the
# normalized CDF of equalized image
img = cv2.imread('imgs/sudoku-original.jpg',0)
W,H = img.shape
img_myeq = myEqualizeHist(img)

# Your implementation may NOT need to return an image that is
# exactly the same as the one OpenCV build-in function does.
# However, the normalized CDF should make sense.
hist_myeq = np.histogram(img_myeq, bins=256, range=(0.0, 255.0))
plt.figure(figsize=(5,5))
plt.plot(range(0,256),np.cumsum(hist_myeq[0])*255/(W*H))
plt.title('Equalized Image: normalized CDF')
plt.show()
```



Threshold

Simple Threshold

If pixel value is greater than a threshold value, it is assigned one value (may be white), else it is assigned another value (may be black). The function used is [cv2.threshold](#).

```
# Get list of available flags for thresholding styles
flags = [i for i in dir(cv2) if i.startswith('THRESH_')]
print(flags)

['THRESH_BINARY', 'THRESH_BINARY_INV', 'THRESH_MASK', 'THRESH_OTSU',
 'THRESH_TOZERO', 'THRESH_TOZERO_INV', 'THRESH_TRIANGLE',
 'THRESH_TRUNC']
```

Adaptive Method

It decides how thresholding value is calculated. The function used is [cv2.adaptiveThreshold](#).

- `cv2.ADAPTIVE_THRESH_MEAN_C` : threshold value is the mean of neighbourhood area.

- `cv2.ADAPTIVE_THRESH_GAUSSIAN_C` : threshold value is the weighted sum of neighbourhood values where weights are a gaussian window.

```
img = cv2.imread('imgs/sudoku-original.jpg',0)
img = cv2.medianBlur(img,5)
img_mean = np.mean(img)

C = 2
ret,th1 = cv2.threshold(img,img_mean,255,cv2.THRESH_BINARY)
th2 = cv2.adaptiveThreshold(img,255,cv2.ADAPTIVE_THRESH_MEAN_C,\
                             cv2.THRESH_BINARY,11,C)

th3 = cv2.adaptiveThreshold(img,255,cv2.ADAPTIVE_THRESH_GAUSSIAN_C,\
                             cv2.THRESH_BINARY,11,C)

#####
#####
# TODO: #
# Trying several value of constant C and observing how the output
#
# thresholded images change.
#
#####
#####
# C_list = [2,4,6,8,10]
# for c in C_list:
#     th2 = cv2.adaptiveThreshold(img,255,cv2.ADAPTIVE_THRESH_MEAN_C,\
#                                 cv2.THRESH_BINARY,11,c)
#
#     th3 =
#     cv2.adaptiveThreshold(img,255,cv2.ADAPTIVE_THRESH_GAUSSIAN_C,\
#                             cv2.THRESH_BINARY,11,c)
#####
#####
#
# END OF YOUR CODE
#
#####
#####

titles = ['Original Image', 'Global Thresholding (v =
{: .2f})'.format(img_mean),
          'Adaptive Mean Thresholding', 'Adaptive Gaussian
Thresholding']
images = [img, th1, th2, th3]


fig = plt.figure(figsize=(10, 10))
for i in range(4):
    plt.subplot(2,2,i+1)
    plt.imshow(images[i],'gray')
    plt.title(titles[i])
```

```
plt.xticks([])
plt.yticks([])
plt.show()
```


Original Image



Global Thresholding (v = 103.69)



Adaptive Mean Thresholding



Adaptive Gaussian Thresholding

